

1 Detailed Theoretical Analysis

Lemma 1.1 (Expected DAC for All-at-Once Fetching). *If the predicted position lies in a page with a uniformly distributed offset, the expected number of I/Os with all-at-once strategy is*

$$\mathbb{E}[DAC] = 1 + \frac{2\varepsilon}{C_{ipp}}, \quad (1)$$

where C_{ipp} represents the number of items per page.

Proof. Denote the offset within the page of the predicted position by $s \sim U(0, C_{ipp} - 1)$. The I/O needs to fetch the page that contains the predicted position and additional pages required to cover the left and right of the $\pm\varepsilon$ window:

$$\mathbb{E}[DAC] = \frac{1}{C_{ipp}} \sum_{s=0}^{C_{ipp}-1} \left(1 + \left\lceil \frac{\varepsilon-s}{C_{ipp}} \right\rceil + \left\lceil \frac{\varepsilon-(C_{ipp}-1-s)}{C_{ipp}} \right\rceil \right). \quad (2)$$

Rewrite $\varepsilon = \lambda \cdot C_{ipp} + r$ for some $\lambda \in \mathbb{N}$ and $0 \leq r < C_{ipp}$. Then we have

$$\begin{aligned} \mathbb{E}[DAC] &= \sum_{s=0}^{C_{ipp}-1} \frac{1}{C_{ipp}} \left(1 + (\lambda + \mathbb{1}_{\{s < r\}}) + (\lambda + \mathbb{1}_{\{s > C_{ipp}-1-r\}}) \right) \\ &= 1 + 2\lambda + \frac{2r}{C_{ipp}} = 1 + \frac{2\varepsilon}{C_{ipp}}. \end{aligned} \quad (3)$$

□

Lemma 1.2. (Expected Cost For One-By-One Fetching). *If the predicted position lies in a page with a uniformly distributed offset, the expected number of I/Os under the one-by-one strategy is*

$$\mathbb{E}[DAC] = 1 + \frac{\varepsilon}{C_{ipp}} \quad (4)$$

Proof. Suppose the predicted position is at \hat{y} and the true position y is uniformly random over the entire search window $[\hat{y} - \varepsilon, \hat{y} + \varepsilon]$. Let X denote the distance between y and the lower bound $\hat{y} - \varepsilon$, i.e., $X \sim U(0, 2\varepsilon)$.

Let k denote a uniformly distributed offset of the lower bound within a page. The expected I/O cost becomes:

$$\mathbb{E}(DAC) = 1 + \frac{1}{2\varepsilon + 1} \sum_{x=0}^{2\varepsilon} \left(\frac{1}{C_{ipp}} \sum_{k=0}^{C_{ipp}-1} \left\lfloor \frac{k+x}{C_{ipp}} \right\rfloor \right). \quad (5)$$

Let $x = q \cdot C_{ipp} + r_1, 0 \leq r_1 < C_{ipp}$,

$$\frac{1}{C_{ipp}} \sum_{k=0}^{C_{ipp}-1} \left\lfloor \frac{k+x}{C_{ipp}} \right\rfloor = q + \frac{r_1}{C_{ipp}}. \quad (6)$$

Let $2\varepsilon = M \cdot C_{ipp} + r_2, 0 \leq r_2 < C_{ipp}$,

$$\begin{aligned} \sum_{x=0}^{2\varepsilon} q &= \sum_{m=0}^{M-1} m \cdot C_{ipp} + M \cdot (r_2 + 1) \\ &= \frac{C_{ipp} \cdot M \cdot (M-1)}{2} + M \cdot (r_2 + 1), \end{aligned} \quad (7)$$

$$\sum_{x=0}^{2\varepsilon} r_1 = M \cdot \frac{C_{\text{ipp}} \cdot (C_{\text{ipp}} - 1)}{2} + \frac{r_2 \cdot (r_2 + 1)}{2}. \quad (8)$$

Combining Equations (5) to (8), we obtain:

$$\begin{aligned} \mathbb{E}[\text{DAC}] &= 1 + \frac{1}{M \cdot C_{\text{ipp}} + r_2 + 1} \left(\frac{(M \cdot C_{\text{ipp}} + r_2)(M \cdot C_{\text{ipp}} + r_2 + 1)}{2 \cdot C_{\text{ipp}}} \right) \\ &= 1 + \frac{M \cdot C_{\text{ipp}} + r_2}{2 \cdot C_{\text{ipp}}} = 1 + \frac{\varepsilon}{C_{\text{ipp}}} \end{aligned} \quad (9)$$

□

Theorem 1.3 (Buffer Hit Rate for Sorted Queries). *Let $\mathcal{K} = (k_1, \dots, k_{|\mathcal{K}|})$ be an array of keys. A query sequence \mathcal{Q} is **sorted** w.r.t. \mathcal{K} if it requests k_i before k_j for all $i < j$. For any such query sequence \mathcal{Q} , suppose the buffer capacity C satisfies*

$$C \geq 1 + \lceil 2\varepsilon/C_{\text{ipp}} \rceil,$$

where C_{ipp} is the number of items per page. Then the cache hit rate equals $h = \frac{R-N}{R}$.

Proof. Let (p_1, \dots, p_R) be the page reference sequence generated by processing the m queries, and partition it by queries:

$$(p_1, \dots, p_R) = \pi_1 \| \pi_2 \| \cdots \| \pi_m,$$

where π_t is the subsequence of page IDs referenced while processing query t . For a learned-index-based engine, query t touches exactly the pages in its last-mile window $W_t = [L_t, H_t]$, with $|W_t| \leq 1 + \lceil \frac{2\varepsilon}{C_{\text{ipp}}} \rceil$. Let \mathcal{P} be the set of distinct pages appearing in (p_1, \dots, p_R) , and $|\mathcal{P}| = N$. Since queries are sorted, the windows move monotonically, i.e., $L_{t+1} \geq L_t$. Hence, between two consecutive queries, only pages newly entering the window can miss: pages in $W_{t+1} \cap W_t$ are still resident and therefore hit, while pages in $W_{t+1} \setminus W_t$ may incur misses. Because $C \geq C_\delta \geq |W_t|$ for all t , the entire window W_t fits in cache during the processing of query t , so no page in W_t can be evicted before π_t finishes. By monotonicity of L_t , once a page is loaded it is either reused by subsequent overlapping windows (and thus hits) or never referenced again. Therefore, each distinct page in \mathcal{P} incurs exactly one compulsory miss—on its first reference—and all later references are hits. The total number of misses is N , so the hit rate is

$$h = \frac{R-N}{R},$$

as claimed. □

Corollary 1.4 (Sorted Order Maximizes Hit Rate). *Given a multiset of queries \mathcal{Q} , the cache hit rate is maximized when \mathcal{Q} is executed in sorted order.*

Proof. For an arbitrary ordering σ , let \mathcal{P} be the set of distinct pages referenced by the resulting trace, and $|\mathcal{P}| = N$. Let $b_\sigma(p)$ denote the number of cache misses (i.e., disk I/Os) incurred by page $p \in \mathcal{P}$. Since each distinct page must be brought into cache at least once, we have $b_\sigma(p) \geq 1$ for all $p \in \mathcal{P}$. Therefore,

$$\text{misses}(\sigma) = \sum_{p \in \mathcal{P}} b_\sigma(p) \geq \sum_{p \in \mathcal{P}} 1 = N = \text{misses(sorted)}. \quad (10)$$


Under sorted order, we have $b_{\text{sorted}}(q) = 1$ for all q and the total number of misses attains this lower bound. It follows that

$$\text{hits}(\sigma) = n - \text{misses}(\sigma) \leq n - N,$$

so the hit rate is maximized under sorted order. □

2 Extended Experimental Evaluation

2.1 Worker Threads Configuration



2.2 Performance of FALCON

2.3 Evaluation of CAM