

# 1 Detailed Theory Analysis

**Lemma 1.1** (Expected DAC for All-at-Once Fetching). *If the predicted position lies in a page with a uniformly distributed offset, the expected number of I/Os with all-at-once strategy is*

$$\mathbb{E}[\text{DAC}] = 1 + \frac{2\varepsilon}{C_{\text{ipp}}}, \quad (1)$$

where  $C_{\text{ipp}}$  represents the number of items per page.

*Proof.* Denote the offset within the page of the predicted position by  $s \sim U(0, C_{\text{ipp}} - 1)$ . The I/O needs to fetch the page that contains the predicted position and additional pages required to cover the left and right of the  $\pm\varepsilon$  window:

$$\mathbb{E}[\text{DAC}] = \frac{1}{C_{\text{ipp}}} \sum_{s=0}^{C_{\text{ipp}}-1} \left( 1 + \left\lceil \frac{\varepsilon-s}{C_{\text{ipp}}} \right\rceil + \left\lceil \frac{\varepsilon-(C_{\text{ipp}}-1-s)}{C_{\text{ipp}}} \right\rceil \right). \quad (2)$$

Rewrite  $\varepsilon = \lambda \cdot C_{\text{ipp}} + r$  for some  $\lambda \in \mathbb{N}$  and  $0 \leq r < C_{\text{ipp}}$ . Then each ceiling term equals  $q$  plus an indicator for crossing a page boundary, yielding

$$\mathbb{E}[\text{DAC}] = 1 + 2\lambda + \frac{2r}{C_{\text{ipp}}} = 1 + \frac{2\varepsilon}{C_{\text{ipp}}}. \quad (3)$$

□

**Lemma 1.2.** (Expected Cost For One-By-One Fetching). *If the predicted position lies in a page with a uniformly distributed offset, the expected number of I/Os under the one-by-one strategy is*

$$\mathbb{E}[\text{DAC}] = 1 + \frac{\varepsilon}{C_{\text{ipp}}} \quad (4)$$

*Proof.* Suppose the predicted position is at  $\hat{y}$  and the true position  $y$  is uniformly random over the entire search window  $[\hat{y} - \varepsilon, \hat{y} + \varepsilon]$ . Let  $X$  denote the distance between  $y$  and the lower bound  $\hat{y} - \varepsilon$ , i.e.,  $X \sim U(0, 2\varepsilon)$ .

Let  $k$  denote a uniformly distributed offset of the lower bound within a page. The expected I/O cost becomes:

$$\mathbb{E}(\text{DAC}) = 1 + \frac{1}{2\varepsilon+1} \sum_{x=0}^{2\varepsilon} \left( \frac{1}{C_{\text{ipp}}} \sum_{k=0}^{C_{\text{ipp}}-1} \left\lfloor \frac{k+x}{C_{\text{ipp}}} \right\rfloor \right). \quad (5)$$

Let  $x = q \cdot C_{\text{ipp}} + r_1, 0 \leq r_1 < C_{\text{ipp}}$ ,

$$\frac{1}{C_{\text{ipp}}} \sum_{k=0}^{C_{\text{ipp}}-1} \left\lfloor \frac{k+x}{C_{\text{ipp}}} \right\rfloor = q + \frac{r_1}{C_{\text{ipp}}}. \quad (6)$$

Let  $2\varepsilon = M \cdot C_{\text{ipp}} + r_2, 0 \leq r_2 < C_{\text{ipp}}$ ,

$$\begin{aligned} \sum_{x=0}^{2\varepsilon} q &= \sum_{m=0}^{M-1} m \cdot C_{\text{ipp}} + M \cdot (r_2 + 1) \\ &= \frac{C_{\text{ipp}} \cdot M \cdot (M-1)}{2} + M \cdot (r_2 + 1), \end{aligned} \quad (7)$$

$$\sum_{x=0}^{2\varepsilon} r_1 = M \cdot \frac{C_{\text{ipp}} \cdot (C_{\text{ipp}}-1)}{2} + \frac{r_2 \cdot (r_2+1)}{2}. \quad (8)$$

Combining Equations (5) to (8), we obtain:

$$\begin{aligned}\mathbb{E}[\text{DAC}] &= 1 + \frac{1}{M \cdot C_{\text{ipp}} + r_2 + 1} \left( \frac{(M \cdot C_{\text{ipp}} + r_2)(M \cdot C_{\text{ipp}} + r_2 + 1)}{2 \cdot C_{\text{ipp}}} \right) \\ &= 1 + \frac{M \cdot C_{\text{ipp}} + r_2}{2 \cdot C_{\text{ipp}}} = 1 + \frac{\varepsilon}{C_{\text{ipp}}}\end{aligned}\tag{9}$$

□

**Theorem 1.3** (Buffer Hit Rate for Sorted Queries). *Let  $\mathcal{K} = (k_1, \dots, k_{|\mathcal{K}|})$  be an array of keys. A query sequence  $\mathcal{Q}$  is **sorted** w.r.t.  $\mathcal{K}$  if it requests  $k_i$  before  $k_j$  for all  $i < j$ . For any such query sequence  $\mathcal{Q}$ , suppose the buffer capacity  $C$  satisfies*

$$C \geq 1 + \lceil 2\varepsilon/C_{\text{ipp}} \rceil,$$

where  $C_{\text{ipp}}$  is the number of items per page. Then the cache hit rate equals  $h = \frac{R-N}{R}$ .

*Proof.* Let  $(p_1, \dots, p_R)$  be the page reference sequence generated by processing the  $m$  queries, and partition it by queries:

$$(p_1, \dots, p_R) = \pi_1 \| \pi_2 \| \cdots \| \pi_m,$$

where  $\pi_t$  is the subsequence of page IDs referenced while processing query  $t$ . For a learned-index-based engine, query  $t$  touches exactly the pages in its last-mile window  $W_t = [L_t, H_t]$ , with  $|W_t| \leq 1 + \lceil \frac{2\varepsilon}{C_{\text{ipp}}} \rceil$ . Let  $\mathcal{P}$  be the set of distinct pages appearing in  $(p_1, \dots, p_R)$ , and  $|\mathcal{P}| = N$ . Since queries are sorted, the windows move monotonically, i.e.,  $L_{t+1} \geq L_t$ . Hence, between two consecutive queries, only pages newly entering the window can miss: pages in  $W_{t+1} \cap W_t$  are still resident and therefore hit, while pages in  $W_{t+1} \setminus W_t$  may incur misses. Because  $C \geq C_\delta \geq |W_t|$  for all  $t$ , the entire window  $W_t$  fits in cache during the processing of query  $t$ , so no page in  $W_t$  can be evicted before  $\pi_t$  finishes. By monotonicity of  $L_t$ , once a page is loaded it is either reused by subsequent overlapping windows (and thus hits) or never referenced again. Therefore, each distinct page in  $\mathcal{P}$  incurs exactly one compulsory miss—on its first reference—and all later references are hits. The total number of misses is  $N$ , so the hit rate is

$$h = \frac{R - N}{R},$$

as claimed. □

**Corollary 1.4** (Sorted Order Maximizes Hit Rate). *Given a multiset of queries  $\mathcal{Q}$ , the cache hit rate is maximized when  $\mathcal{Q}$  is executed in sorted order.*

*Proof.* For an arbitrary ordering  $\sigma$ , let  $\mathcal{P}$  be the set of distinct pages referenced by the resulting trace, and  $|\mathcal{P}| = N$ . Let  $b_\sigma(p)$  denote the number of cache misses (i.e., disk I/Os) incurred by page  $p \in \mathcal{P}$ . Since each distinct page must be brought into cache at least once, we have  $b_\sigma(p) \geq 1$  for all  $p \in \mathcal{P}$ . Therefore,

$$\text{misses}(\sigma) = \sum_{p \in \mathcal{P}} b_\sigma(p) \geq \sum_{p \in \mathcal{P}} 1 = N = \text{misses(sorted)}.\tag{10}$$

Under sorted order, we have  $b_{\text{sorted}}(q) = 1$  for all  $q$  and the total number of misses attains this lower bound. It follows that


$$\text{hits}(\sigma) = n - \text{misses}(\sigma) \leq n - N,$$

so the hit rate is maximized under sorted order. □


## 2 Extended Experiments

### 2.1 Worker Threads Configuration

In this section, we exhibit Point-query performance of FALCON under varying numbers of producer threads ( $P$ ) and worker threads ( $W$ ). The experiments are conducted on books dataset.



(a) Throughput (KOPS) without page buffer.



(b) Latency (ms) without page buffer.

### 2.2 Performance of FALCON

### 2.3 Evaluation of CAM